

Chris Jordan
Victor Janmey
Raymond Ko

PARALLELIZING K-MEANS FOR IMAGE SEGMENTATION

Problems and Applications

- ⦿ Group points to clusters
- ⦿ Fixed number of clusters (k)
- ⦿ Machine learning
 - Unsupervised learning methods
 - Dividing data for more specific classifiers
 - Selection of examples for k-nearest neighbors
- ⦿ Computer vision
 - Image segmentation
 - Estimating object boundaries
 - Cue for object recognition
- ⦿ Network Analysis
 - Finding communities



K-means Algorithm

⦿ Step 1: Initialization

- Choose initial centers
 - Naïve approach
 - K-means++ (Significant speedup)

⦿ Step 2: Iteration

- Assign particle to centers (distance measure)
- Re-compute cluster centers (means of points)
- Terminate when assignments converge

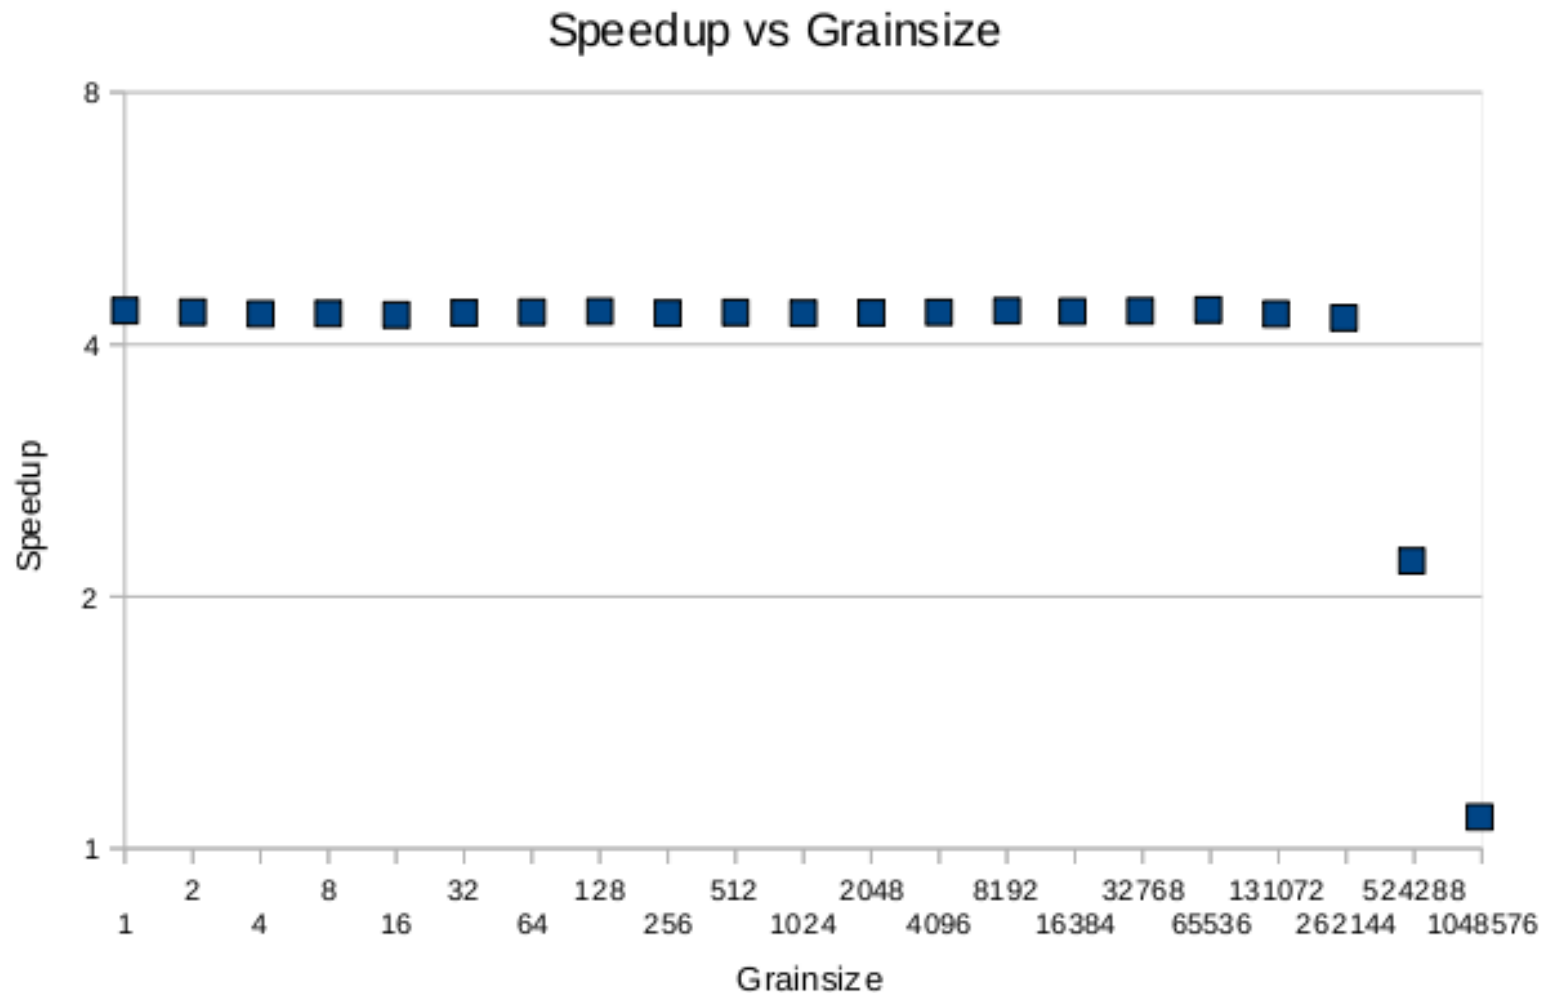
Implementation

- ◎ C++ and OpenCV image library
 - Required for TBB, OpenCV has easy image manipulation, input, and output
- ◎ Parallel: Add TBB constructs to serial
- ◎ Parallel cluster assignment: `parallel_for`
- ◎ Parallel cluster center calculation:
`parallel_reduce`

Preliminary results

- Sample run: 1,024,000 pixels, 3 dimensions (HSV)
- 50 Iteration, 50 clusters
- Speedup: 4.4, Grain Size: 64,000
- Step 1 speedup: 4.42, grain size: 64,000
- Step 2 speedup: 2.89 grain size: 8,192
- Step 1 dominates since it is $(k*n)$ vs. $O(n)$ for Step 2
- Similarly for k-means++ initial cluster selection, Step 1 (min distance to a cluster) is the dominating factor since it is also $O(n)$

Speedup vs. Grainsize



Next steps

- ⦿ Investigate if similar performance is achievable with OpenMP
- ⦿ Experiments
 - Speedup vs. Number of Particles
 - Is there an optimal image size given today's technology that balances speed and benefits?
 - Speedup vs. Number of Cores
 - Does it scale well? When will performance gains be apparent?
 - Speedup vs. Number of Clusters
 - At what point does it become inefficient to use the algorithm, if any?
- ⦿ Algorithm improvements
 - K-means++
 - Triangle inequality