

Parallelizing K-means for Image Segmentation

Chris Jordan, Victor Janmey, and
Raymond Ko

Problem and Applications

- Group points into k clusters
- Machine Learning
 - Unsupervised learning methods
 - Dividing data for more specific classifiers
 - Selection of examples for k -nearest neighbors
- Computer Vision
 - Image segmentation
 - Estimating object boundaries
 - Cue for object recognition



K-means Algorithm

- Initialization:
 - Select initial centers
- Iteration:
 - Step 1:
Assign each particle to nearest cluster
 - Step 2:
Recompute cluster centers
 - Repeat until convergence
- Step 1 dominates since it is $(k*n)$ vs. $O(n)$ for step 2




Implementation

- C++ and OpenCV image library
- Parallel implementation with TBB
- Cluster assignment: `parallel_for`
- Cluster center calculation: `parallel_reduce`

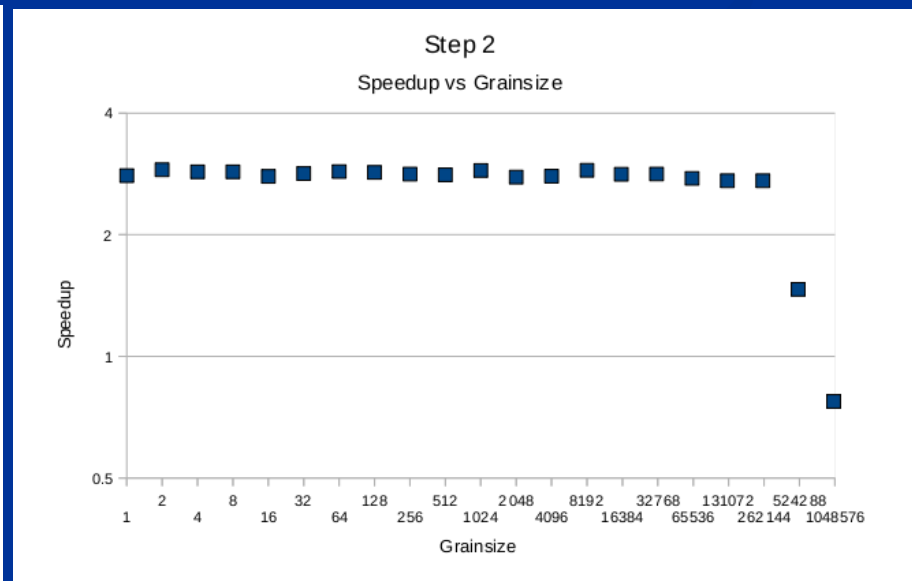
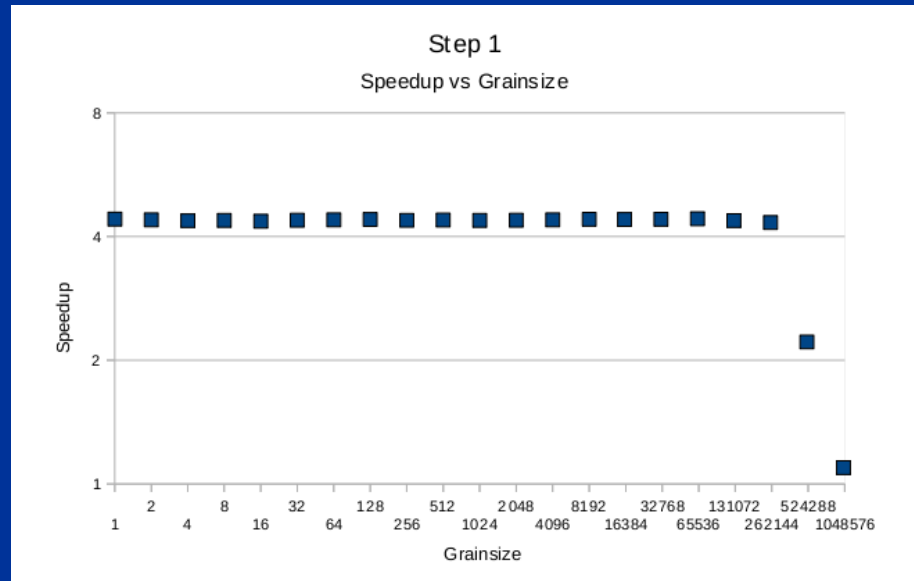
Preliminary Results

- Sample run with 50 clusters on 4 cores
- 1,024,000 pixels over 3 dimensions (HSV)
- 4.4 speedup at Grain Size: 64,000
- Superlinear because of loop unrolling in TBB?



Results by Step

- Step 1 speedup: 4.42 at grain size of 64,000
- Step 2 speedup: 2.89 at grain size of 8,192



Next Steps

- OpenMP
- Experiments
 - Speedup vs. Number of Particles
 - Speedup vs. Number of Cores
 - Speedup vs. Number of Clusters
- Algorithm improvements
 - K-means++
 - Triangle inequality